

Introduction to PyTorch

Nikos Gkanatsios

Credits to Sai Shruthi Balaji

Preliminaries: NumPy

```
In [1]: import numpy as np
```

```
In [2]: a = np.array([1,2,3,4,5,6,7,8,9])
```

```
In [3]: a
```

```
Out[3]: array([1, 2, 3, 4, 5, 6, 7, 8, 9])
```


```
In [4]: b = a.reshape((3,3))
```

```
In [5]: b
```

```
Out[5]:  
array([[1, 2, 3],  
       [4, 5, 6],  
       [7, 8, 9]])
```

```
In [6]: b * 10 + 4
```

```
Out[6]:  
array([[14, 24, 34],  
       [44, 54, 64],  
       [74, 84, 94]])
```




Expression

`arr[:2, 1:]`

Shape

`(2, 2)`



`arr[2]`


`(3,)`

`arr[2, :]`

`(3,)`


`arr[2:, :]`

`(1, 3)`



`arr[:, :2]`

`(3, 2)`



`arr[1, :2]`

`(2,)`

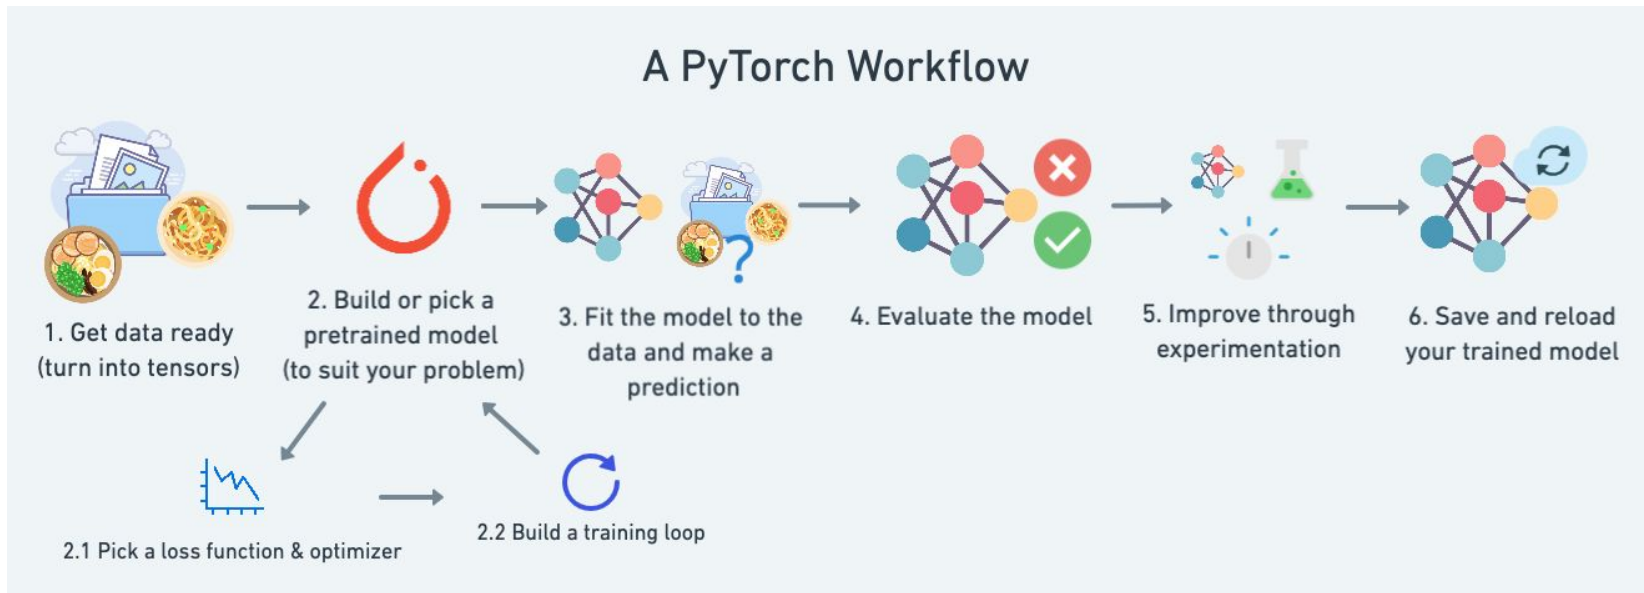
`arr[1:2, :2]`

`(1, 2)`

What is PyTorch?




- An open-source machine learning framework that accelerates the path from research prototyping to production deployment.



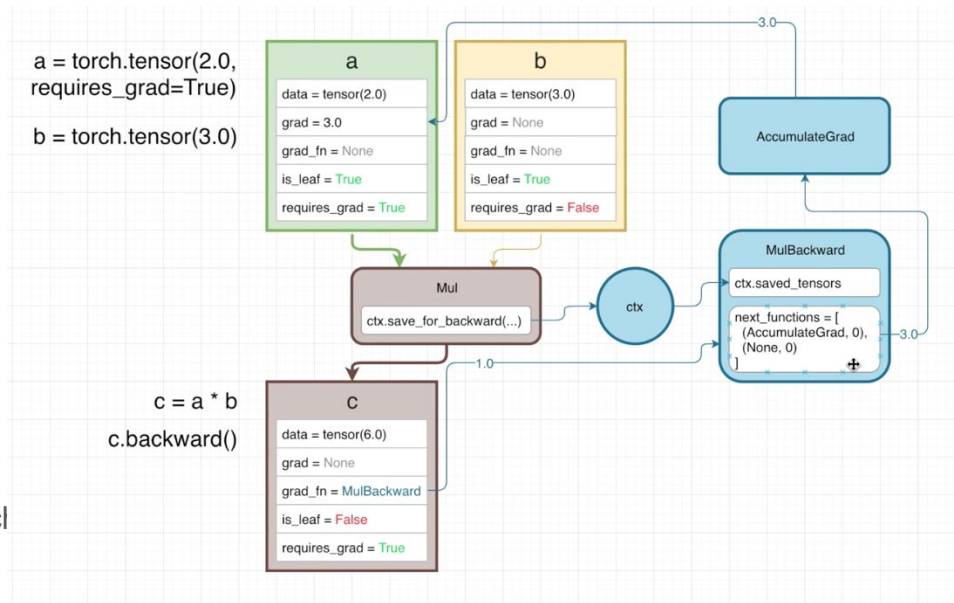
Tensors: Backbone of PyTorch

- Multi-dimensional array, same as numpy array
- Biggest difference: Tensors can be run on CPU/GPU



Tensors Operations

1. Directly stores input data, weights, etc
2. Holds grad, grad_fn (requires_grad)
3. detach()
4. clone()
5. numpy()
6. tensor.device
7. to(device) -> CPU or GPU / which GPU? (cuda:0)
8. tensor.dtype
9. to(dtype)
10. Between Tensors: similar operations like numpy such as:
 - Add, Matmul, Subtract, Concat, etc.

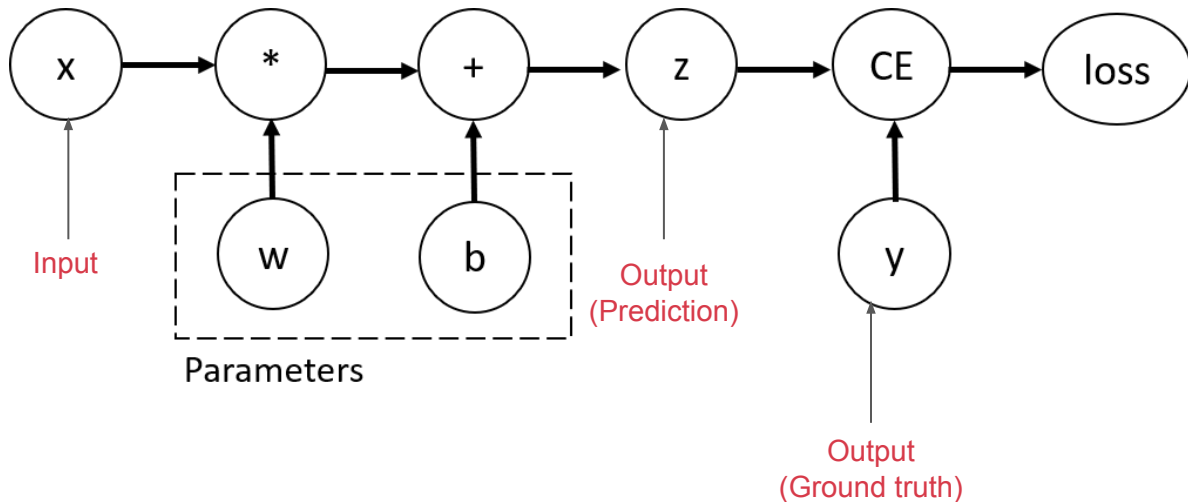


Autograd

- This class is PyTorch's automatic differentiation engine that powers neural network training.

```
import torch

x = torch.ones(5) # input tensor
y = torch.zeros(3) # expected output
w = torch.randn(5, 3, requires_grad=True)
b = torch.randn(3, requires_grad=True)
z = torch.matmul(x, w)+b
loss = torch.nn.functional.binary_cross_entropy_with_logits(z, y)
```




Autograd

```
loss.backward()  
print(w.grad)  
print(b.grad)
```

Out:

```
tensor([[0.0659, 0.0797, 0.2611],  
        [0.0659, 0.0797, 0.2611],  
        [0.0659, 0.0797, 0.2611],  
        [0.0659, 0.0797, 0.2611],  
        [0.0659, 0.0797, 0.2611]])  
tensor([0.0659, 0.0797, 0.2611])
```




Parts of a Framework

- Data loading
- Model definition: Forward Pass
- Loss
- Backward Pass
- Optimization
- Validation

Data Loading

- Load data, turn into tensors, and batch
- Common image datasets available in torchvision.datasets
- DataLoader can handle shuffling and batching
- Writing your own DataSet class:
 - `__init__`
 - `__len__`
 - `__getitem__`



Defining a model: nn.Module

- Model class extends nn.Module and has the following methods:
 - `__init__`
 - `forward`
- nn.Module advantages:
 - Module reuse
 - Easy chaining of multiple steps
 - Intermediate states are held in compute graph


```
class TinyModel(torch.nn.Module):  
  
    def __init__(self):  
        super(TinyModel, self).__init__()  
  
        self.linear1 = torch.nn.Linear(100, 200)  
        self.activation = torch.nn.ReLU()  
        self.linear2 = torch.nn.Linear(200, 10)  
        self.softmax = torch.nn.Softmax()  
  
    def forward(self, x):  
        x = self.linear1(x)  
        x = self.activation(x)  
        x = self.linear2(x)  
        x = self.softmax(x)  
        return x
```

nn.Module: Useful Facts

1. `module.parameters()` -> gathers all `nn.Parameter` in the module
2. `parameters()` -> parameters to “update/optimize”
3. `state_dict()` -> dictionary of all parameters and buffers -> `torch.save`
4. `load_state_dict(state_dict)` -> used to load pretrained weights
5. `train()` - trainable weights
6. `eval()` - frozen dropout/norm layers
7. `to(device)` - train on GPU

Loss function

- Measures how well the model is doing
- Key for training:
 - Used to compute gradients
 - Used in back propagation (`loss.backward`)
- Different loss functions for different problems:
 - Regression: L1 Loss, L2 Loss
 - Classification: Cross Entropy Loss, NLL Loss, BCE Loss
 - Ranking: Margin Ranking Loss, Triplet Margin Loss
 - KL Divergence Loss: Difference between distributions



Optimization: torch.optim


- Pass parameters, define learning rate and other optional params.

```
optim.SGD([  
    {'params': model.base.parameters()},  
    {'params': model.classifier.parameters(), 'lr': 1e-3}  
], lr=1e-2, momentum=0.9)
```

- Set gradients as zero and take a step.

```
for input, target in dataset:  
    optimizer.zero_grad()  
    output = model(input)  
    loss = loss_fn(output, target)  
    loss.backward()  
    optimizer.step()
```

- Common optimizers: Adam, SGD, RMSProp



Validation

- Download test dataset, run it periodically to look at accuracy/loss
- Set `model.eval()` to deactivate dropout and norm layers from updating
- Run with `torch.no_grad` to deactivate gradients
- Use an evaluation metric.
 - Eg: mAP for classification
 - Eg: IOU for bounding box prediction

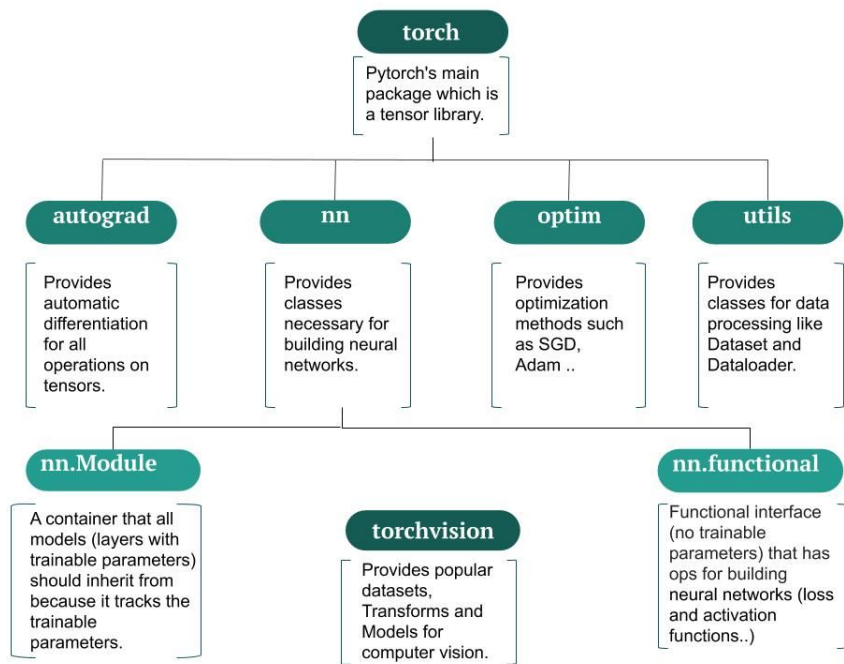
Running on the GPU

<https://pytorch.org/docs/stable/notes/cuda.html>

1. Move your dataset and your model to the GPU
2. Simple: `.cuda()`
3. Can also create the device and do `.to(device)`

Summary

1. Load data
 - a. torchvision is sometimes convenient
 - b. DataLoader for batching, shuffling
2. Define net
 - a. Convenient to implement nn.Module, to get parameters, zero_grad, etc easily
3. Define loss
 - a. Pytorch provides a lot of these, eg in torch.nn.functional
4. Backward automatically computes gradients
 - a. Can manually clip, etc if desired
5. Optimizer to update the given parameters
 - a. torch.optim



References

- <https://pytorch.org/tutorials/beginner/blitz/>
- <https://pytorch.org/tutorials/beginner/basics/intro.html>
- https://pytorch.org/tutorials/beginner/basics/autogradqs_tutorial.html
- <https://github.com/yunjey/pytorch-tutorial>
- <https://pytorch.org/tutorials/beginner/ptcheat.html>
- https://web.cs.ucdavis.edu/~yjee/teaching/ecs289g-winter2018/Pytorch_Tutorial.pdf
- <https://www.kaggle.com/code/residentmario/pytorch-autograd-explained/notebook>
- <https://neptune.ai/blog/pytorch-loss-functions>